

Data wrangling

El Mex

Contents

1 Basic functions from dplyr.	1
1.1 filter rows	2
1.2 summarize variables	4
1.3 group_by rows	4
1.4 mutate existing variables	6
1.5 arrange and sort rows	10
1.6 join data frames	11
1.7 select variables/columns	15
1.8 rename variables	18
1.9 top_n values of a variable	18

1 Basic functions from dplyr.

Needed packages

```
library(dplyr)
library(ggplot2)
library(nycflights13)
```

```
# explore the dataframe called 'flights'
glimpse(flights)
```

```
## Rows: 336,776
## Columns: 19
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013...
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 55...
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 60...
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2,...
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 8...
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 8...
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7,...
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6"...
```

```
## $ flight      <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301...
## $ tailnum     <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N...
## $ origin      <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LG...
## $ dest        <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IA...
## $ air_time    <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149...
## $ distance    <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 73...
## $ hour        <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6...
## $ minute      <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59...
## $ time_hour   <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-0...
```

1.1 filter rows

```
# flights from New York City to Portland, Oregon (destination code: "PDX")
portland_flights <- flights %>%
  filter(dest == "PDX")

# see the first 6 rows
head(portland_flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1    1739           1740        -1     2051           2112
## 2  2013     1     1    1805           1757         8     2117           2119
## 3  2013     1     1    2052           2029        23     2349           2350
## 4  2013     1     2     804            805        -1     1039           1110
## 5  2013     1     2    1552           1550         2     1853           1922
## 6  2013     1     2    1727           1720         7     2042           2040
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Let's begin with an exercise

```
# filter all rows from JFK that were heading to Burlington ("BTV") or Seattle ("SEA") &
# in the months of October, November, or December.
btv_sea_flights_fall <- flights %>%
  filter(origin == "JFK" & (dest == "BTV" | dest == "SEA") & month >= 10)

# see the first 6 rows
head(btv_sea_flights_fall)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013    10     1     729            735        -6     1049           1040
## 2  2013    10     1     853            900        -7     1217           1157
## 3  2013    10     1     916            925        -9     1016           1033
## 4  2013    10     1    1216           1221        -5     1326           1328
## 5  2013    10     1    1452           1459        -7     1602           1622
```

```
## 6 2013 10 1 1459 1500 -1 1817 1829
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

We could use a comma instead of “&” to get the same results

```
btv_sea_flights_fall <- flights %>%
  filter(origin == "JFK", (dest == "BTV" | dest == "SEA"), month >= 10)
```

Using ! “not” operator to pick rows that don’t match a criteria

```
# filtering rows to flights that didn't go to Burlington, "BTV" or Seattle, "SEA"
not_BTV_SEA <- flights %>%
  filter(!(dest == "BTV" | dest == "SEA"))

head(not_BTV_SEA)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     517             515           2     830           819
## 2 2013     1     1     533             529           4     850           830
## 3 2013     1     1     542             540           2     923           850
## 4 2013     1     1     544             545          -1    1004          1022
## 5 2013     1     1     554             600          -6     812           837
## 6 2013     1     1     554             558          -4     740           728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

Warning: note the parentheses around the (dest == “BTV” | dest == “SEA”). Let’s say we put it this way:

```
flights %>% filter(!dest == "BTV" | dest == "SEA")
```

This would give us all flights not headed to “BTV” or those headed to “SEA”, clearly a different result.

Now say we have a larger number of airports we want to filter for, say “SEA”, “SFO”, “PDX”, “BTV”, and “BDL”. Using the | (or) operator isn’t quite practical

```
many_airports <- flights %>%
  filter(dest == "SEA" | dest == "SFO" | dest == "PDX" |
         dest == "BTV" | dest == "BDL")
```

A better idea is to use the %in% operator

```
many_airports <- flights %>%
  filter(dest %in% c("SEA", "SFO", "PDX", "BTV", "BDL"))

head(many_airports)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     558             600        -2      923             937
## 2  2013     1     1     611             600         11      945             931
## 3  2013     1     1     655             700        -5     1037            1045
## 4  2013     1     1     724             725        -1     1020            1030
## 5  2013     1     1     729             730        -1     1049            1115
## 6  2013     1     1     734             737        -3     1047            1113
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

1.2 summarize variables

Calculate summary stats of the temp variable in the weather data frame

```
# create variable 'summary_temp' & get both, mean and standard deviation of 'temp'
summary_temp <- weather %>%
  summarize(mean = mean(temp), std_dev = sd(temp))

# check it out
head(summary_temp)
```

```
## # A tibble: 1 x 2
##   mean std_dev
##   <dbl>   <dbl>
## 1    NA      NA
```

Why did we get both NA as result? it's because we have some NA (missing values) on the the dataset. Let's use na.rm = TRUE to ignore any NA

```
summary_temp <- weather %>%
  summarize(mean = mean(temp, na.rm = TRUE),
            std_dev = sd(temp, na.rm = TRUE))

# check it out
head(summary_temp)
```

```
## # A tibble: 1 x 2
##   mean std_dev
##   <dbl>   <dbl>
## 1  55.3    17.8
```

1.3 group_by rows

Instead of a single mean temperature for the whole year, you would like one for each of the 12 months separately

```
summary_monthly_temp <- weather %>%
  group_by(month) %>%
  summarize(mean = mean(temp, na.rm = TRUE),
            std_dev = sd(temp, na.rm = TRUE))

#check it out
summary_monthly_temp
```

```
## # A tibble: 12 x 3
##   month mean std_dev
##   <int> <dbl> <dbl>
## 1     1  35.6  10.2
## 2     2  34.3   6.98
## 3     3  39.9   6.25
## 4     4  51.7   8.79
## 5     5  61.8   9.68
## 6     6  72.2   7.55
## 7     7  80.1   7.12
## 8     8  74.5   5.19
## 9     9  67.4   8.47
## 10    10  60.1   8.85
## 11    11  45.0  10.4
## 12    12  38.4   9.98
```

Let's use the `n()` counting summary function (it counts rows). Suppose we'd like to count how many flights departed each of the three airports in New York City

```
by_origin <- flights %>%
  group_by(origin) %>%
  summarize(count = n())

#check it out
by_origin
```

```
## # A tibble: 3 x 2
##   origin count
##   <chr>   <int>
## 1 EWR    120835
## 2 JFK    111279
## 3 LGA    104662
```

Say you want to know the number of flights leaving each of the three New York City airports for each month (grouping by more than one variable)

```
by_origin_monthly <- flights %>%
  group_by(origin, month) %>%
  summarize(count = n())

# check it out
head(by_origin_monthly)
```

```
## # A tibble: 6 x 3
## # Groups:   origin [1]
##   origin month count
##   <chr>   <int> <int>
## 1 EWR      1  9893
## 2 EWR      2  9107
## 3 EWR      3 10420
## 4 EWR      4 10531
## 5 EWR      5 10592
## 6 EWR      6 10175
```

When grouping by more than two variables, remember to include all variables at the same time in the same `group_by()`. Otherwise, look at this:

```
by_origin_monthly_incorrect <- flights %>%
  group_by(origin) %>%
  group_by(month) %>%
  summarize(count = n())

# check it out
by_origin_monthly_incorrect
```

```
## # A tibble: 12 x 2
##   month count
##   <int> <int>
## 1     1  27004
## 2     2  24951
## 3     3  28834
## 4     4  28330
## 5     5  28796
## 6     6  28243
## 7     7  29425
## 8     8  29327
## 9     9  27574
## 10    10  28889
## 11    11  27268
## 12    12  28135
```

Here `group_by(month)` overwrote the grouping structure meta-data of the earlier `group_by(origin)`

1.4 mutate existing variables

Convert temperatures from °F to °C with the formula: $\text{temp in } ^\circ\text{C} = (\text{temp in } ^\circ\text{F} - 32) / 1.8$

```
# use 'temp' variable for calculations
weather <- weather %>%
  mutate(temp_in_C = (temp - 32) / 1.8)

# check it out
head(weather)
```

```
## # A tibble: 6 x 16
##   origin year month   day hour temp dewp humid wind_dir wind_speed wind_gust
##   <chr>  <int> <int> <int> <int> <dbl> <dbl> <dbl>   <dbl>    <dbl>    <dbl>
## 1 EWR    2013     1     1     1 39.0 26.1 59.4     270     10.4      NA
## 2 EWR    2013     1     1     2 39.0 27.0 61.6     250      8.06     NA
## 3 EWR    2013     1     1     3 39.0 28.0 64.4     240     11.5      NA
## 4 EWR    2013     1     1     4 39.9 28.0 62.2     250     12.7      NA
## 5 EWR    2013     1     1     5 39.0 28.0 64.4     260     12.7      NA
## 6 EWR    2013     1     1     6 37.9 28.0 67.2     240     11.5      NA
## # ... with 5 more variables: precip <dbl>, pressure <dbl>, visib <dbl>,
## #   time_hour <dtm>, temp_in_C <dbl>
```

Let's now compute monthly average temperatures in both °F and °C using the `group_by()` and `summarize()`

```
summary_monthly_temp <- weather %>%
  group_by(month) %>%
  summarize(mean_temp_in_F = mean(temp, na.rm = TRUE),
            mean_temp_in_C = mean(temp_in_C, na.rm = TRUE))

# check it out
summary_monthly_temp
```

```
## # A tibble: 12 x 3
##   month mean_temp_in_F mean_temp_in_C
##   <int>         <dbl>         <dbl>
## 1     1          35.6           2.02
## 2     2          34.3           1.26
## 3     3          39.9           4.38
## 4     4          51.7          11.0
## 5     5          61.8          16.6
## 6     6          72.2          22.3
## 7     7          80.1          26.7
## 8     8          74.5          23.6
## 9     9          67.4          19.7
## 10    10          60.1          15.6
## 11    11          45.0           7.22
## 12    12          38.4           3.58
```

Passengers are frustrated when their flight departs late, but aren't as annoyed if pilots can make up some time during the flight. This is known in the airline industry as gain. Let's create the variable "gain"

```
flights <- flights %>%
  mutate(gain = dep_delay - arr_delay)

# check it out
head(flights)
```

```
## # A tibble: 6 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
```

```
## 1 2013 1 1 517 515 2 830 819
## 2 2013 1 1 533 529 4 850 830
## 3 2013 1 1 542 540 2 923 850
## 4 2013 1 1 544 545 -1 1004 1022
## 5 2013 1 1 554 600 -6 812 837
## 6 2013 1 1 554 558 -4 740 728
## # ... with 12 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, gain <dbl>
```

The flight in the first row departed 2 minutes late but arrived 11 minutes late, so its “gained time in the air” is a loss of 9 minutes, hence its gain is $2 - 11 = -9$. On the other hand, the flight in the fourth row departed a minute early (dep_delay of -1) but arrived 18 minutes early (arr_delay of -18), so its “gained time in the air” is 17 minutes, hence its gain is +17

Let’s look at some summary stats of gain

```
gain_summary <- flights %>%
  summarize(
    min = min(gain, na.rm = TRUE),
    q1 = quantile(gain, 0.25, na.rm = TRUE),
    median = quantile(gain, 0.5, na.rm = TRUE),
    q3 = quantile(gain, 0.75, na.rm = TRUE),
    max = max(gain, na.rm = TRUE),
    mean = mean(gain, na.rm = TRUE),
    sd = sd(gain, na.rm = TRUE),
    missing = sum(is.na(gain))
  )

gain_summary
```


```
## # A tibble: 1 x 8
##   min    q1 median    q3    max  mean    sd missing
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <int>
## 1 -196    -3     7    17   109  5.66  18.0   9430
```

This code was a bit long to type, we’ll see a succinct way to do it by using `skim()` from `skimr` package later on

We can visualize gain as it is a numerical variable. Let’s make an histogram

```
ggplot(data = flights, mapping = aes(x = gain)) +
  geom_histogram(color = "white", bins = 20)
```

```
## Warning: Removed 9430 rows containing non-finite values (stat_bin).
```

Finally, we can create multiple new variables at once in the same `mutate()` code

```
flights <- flights %>%
  mutate(
    gain = dep_delay - arr_delay,
    hours = air_time / 60,
    gain_per_hour = gain / hours
  )

head(flights)
```

```
## # A tibble: 6 x 22
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>
## 1  2013     1     1     517             515         2        830           819
## 2  2013     1     1     533             529         4        850           830
## 3  2013     1     1     542             540         2        923           850
## 4  2013     1     1     544             545        -1       1004          1022
## 5  2013     1     1     554             600        -6        812           837
## 6  2013     1     1     554             558        -4        740           728
## # ... with 14 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, gain <dbl>, hours <dbl>,
## #   gain_per_hour <dbl>
```

1.5 arrange and sort rows

Suppose we are interested in determining the most frequent destination airports for all domestic flights departing from New York City in 2013

```
freq_dest <- flights %>%  
  group_by(dest) %>%  
  summarize(num_flights = n())
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
#check it out  
head(freq_dest)
```

```
## # A tibble: 6 x 2  
##   dest  num_flights  
##   <chr>      <int>  
## 1 ABQ         254  
## 2 ACK         265  
## 3 ALB         439  
## 4 ANC          8  
## 5 ATL       17215  
## 6 AUS       2439
```

Observe that by default the rows are sorted in alphabetical order. Say instead we would like to see it sorted from the most to the least number of flights (`num_flights`) instead

```
freq_dest_default <- freq_dest %>%  
  arrange(num_flights)
```

```
# check it out  
head(freq_dest_default, 10)
```

```
## # A tibble: 10 x 2  
##   dest  num_flights  
##   <chr>      <int>  
## 1 LEX          1  
## 2 LGA          1  
## 3 ANC          8  
## 4 SBN         10  
## 5 HDN         15  
## 6 MTJ         15  
## 7 EYW         17  
## 8 PSP         19  
## 9 JAC         25  
## 10 BZN        36
```

As we see, “ascending” order is default. To switch the ordering to be in “descending” order, we use the `desc()`

```
freq_dest_desc <- freq_dest %>%
  arrange(desc(num_flights))

head(freq_dest_desc, 10)
```

```
## # A tibble: 10 x 2
##   dest num_flights
##   <chr>      <int>
## 1 ORD      17283
## 2 ATL      17215
## 3 LAX      16174
## 4 BOS      15508
## 5 MCO      14082
## 6 CLT      14064
## 7 SFO      13331
## 8 FLL      12055
## 9 MIA      11728
## 10 DCA       9705
```

1.6 join data frames

In both the `flights` and `airlines` data frames, the key variable we want to join/merge/match the rows by has the same name: `carrier`. Let's use the `inner_join()` function to join the two data frames

```
flights_joined <- flights %>%
  inner_join(airlines, by = "carrier")

# check both data frames
glimpse(flights)
```

```
## Rows: 336,776
## Columns: 22
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013...
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 55...
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 60...
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2,...
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 8...
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 8...
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7,...
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6"...
## $ flight    <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301...
## $ tailnum   <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N...
## $ origin    <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LG...
## $ dest      <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IA...
## $ air_time  <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149...
## $ distance  <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 73...
## $ hour      <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6...
## $ minute    <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59...
## $ time_hour <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-0...
```

```
## $ gain          <dbl> -9, -16, -31, 17, 19, -16, -24, 11, 5, -10, 0, 1, -9...
## $ hours         <dbl> 3.7833333, 3.7833333, 2.6666667, 3.0500000, 1.933333...
## $ gain_per_hour <dbl> -2.3788546, -4.2290749, -11.6250000, 5.5737705, 9.82...
```

```
glimpse(flights_joined)
```

```
## Rows: 336,776
## Columns: 23
## $ year          <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013...
## $ month         <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ day           <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ dep_time      <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 55...
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 60...
## $ dep_delay     <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2,...
## $ arr_time      <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 8...
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 8...
## $ arr_delay     <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7,...
## $ carrier       <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6"...
## $ flight        <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301...
## $ tailnum       <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N...
## $ origin        <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LG...
## $ dest          <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IA...
## $ air_time      <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149...
## $ distance      <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 73...
## $ hour          <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6...
## $ minute        <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59...
## $ time_hour     <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-0...
## $ gain          <dbl> -9, -16, -31, 17, 19, -16, -24, 11, 5, -10, 0, 1, -9...
## $ hours         <dbl> 3.7833333, 3.7833333, 2.6666667, 3.0500000, 1.933333...
## $ gain_per_hour <dbl> -2.3788546, -4.2290749, -11.6250000, 5.5737705, 9.82...
## $ name          <chr> "United Air Lines Inc.", "United Air Lines Inc.", "A...
```

```
head(flights_joined, 10)
```

```
## # A tibble: 10 x 23
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## 7  2013     1     1     555             600          -5     913             854
## 8  2013     1     1     557             600          -3     709             723
## 9  2013     1     1     557             600          -3     838             846
## 10 2013     1     1     558             600          -2     753             745
## # ... with 15 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, gain <dbl>, hours <dbl>,
## #   gain_per_hour <dbl>, name <chr>
```

In airports the airport code is in “faa”, whereas in flights the airport codes are in “origin” and “dest”. In order to join these two data frames by airport code, our inner_join() operation will use the by = c("dest" = "faa")

```
flights_with_airport_names <- flights %>%
  inner_join(airports, by = c("dest" = "faa"))

# explore it
glimpse(flights_with_airport_names)
```

```
## Rows: 329,174
## Columns: 29
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013...
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ dep_time  <int> 517, 533, 542, 554, 554, 555, 557, 557, 558, 558, 55...
## $ sched_dep_time <int> 515, 529, 540, 600, 558, 600, 600, 600, 600, 600, 60...
## $ dep_delay <dbl> 2, 4, 2, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2, -1,...
## $ arr_time  <int> 830, 850, 923, 812, 740, 913, 709, 838, 753, 849, 85...
## $ sched_arr_time <int> 819, 830, 850, 837, 728, 854, 723, 846, 745, 851, 85...
## $ arr_delay <dbl> 11, 20, 33, -25, 12, 19, -14, -8, 8, -2, -3, 7, -14,...
## $ carrier   <chr> "UA", "UA", "AA", "DL", "UA", "B6", "EV", "B6", "AA"...
## $ flight    <int> 1545, 1714, 1141, 461, 1696, 507, 5708, 79, 301, 49,...
## $ tailnum   <chr> "N14228", "N24211", "N619AA", "N668DN", "N39463", "N...
## $ origin    <chr> "EWR", "LGA", "JFK", "LGA", "EWR", "EWR", "LGA", "JF...
## $ dest      <chr> "IAH", "IAH", "MIA", "ATL", "ORD", "FLL", "IAD", "MC...
## $ air_time  <dbl> 227, 227, 160, 116, 150, 158, 53, 140, 138, 149, 158...
## $ distance  <dbl> 1400, 1416, 1089, 762, 719, 1065, 229, 944, 733, 102...
## $ hour      <dbl> 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6, 6...
## $ minute    <dbl> 15, 29, 40, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59, 0,...
## $ time_hour <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-0...
## $ gain      <dbl> -9, -16, -31, 19, -16, -24, 11, 5, -10, 0, 1, -9, 12...
## $ hours     <dbl> 3.7833333, 3.7833333, 2.6666667, 1.9333333, 2.500000...
## $ gain_per_hour <dbl> -2.3788546, -4.2290749, -11.6250000, 9.8275862, -6.4...
## $ name      <chr> "George Bush Intercontinental", "George Bush Interco...
## $ lat       <dbl> 29.98443, 29.98443, 25.79325, 33.63672, 41.97860, 26...
## $ lon       <dbl> -95.34144, -95.34144, -80.29056, -84.42807, -87.9048...
## $ alt       <dbl> 97, 97, 8, 1026, 668, 9, 313, 96, 668, 19, 26, 126, ...
## $ tz        <dbl> -6, -6, -5, -5, -6, -5, -5, -5, -6, -5, -5, -8, -8, ...
## $ dst       <chr> "A", "A", "A", "A", "A", "A", "A", "A", "A", "A", "A...
## $ tzone     <chr> "America/Chicago", "America/Chicago", "America/New_Y..."
```

```
head(flights_with_airport_names)
```

```
## # A tibble: 6 x 29
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>     <int>         <int>
## 1  2013     1     1     517             515           2       830             819
## 2  2013     1     1     533             529           4       850             830
## 3  2013     1     1     542             540           2       923             850
## 4  2013     1     1     554             600          -6       812             837
## 5  2013     1     1     554             558          -4       740             728
## 6  2013     1     1     555             600          -5       913             854
```

```
## # ... with 21 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, gain <dbl>, hours <dbl>,
## #   gain_per_hour <dbl>, name <chr>, lat <dbl>, lon <dbl>, alt <dbl>, tz <dbl>,
## #   dst <chr>, tzone <chr>
```

Let's construct the chain of pipe operators `%>%` that computes the number of flights from NYC to each destination, but also includes information about each destination airport

```
named_dests <- flights %>%
  group_by(dest) %>%
  summarize(num_flights = n()) %>%
  arrange(desc(num_flights)) %>%
  inner_join(airports, by = c("dest" = "faa")) %>%
  rename(airport_name = name)
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
#explore it
head(named_dests)
```

```
## # A tibble: 6 x 9
##   dest num_flights airport_name      lat   lon alt   tz dst tzone
##   <chr>      <int> <chr>          <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1 ORD         17283 Chicago Ohare Intl  42.0  -87.9  668   -6 A America/~
## 2 ATL         17215 Hartsfield Jackson~  33.6  -84.4  1026  -5 A America/~
## 3 LAX         16174 Los Angeles Intl  33.9 -118.   126   -8 A America/~
## 4 BOS         15508 General Edward Law~  42.4  -71.0   19   -5 A America/~
## 5 MCO         14082 Orlando Intl    28.4  -81.3   96   -5 A America/~
## 6 CLT         14064 Charlotte Douglas ~  35.2  -80.9   748  -5 A America/~
```

Say instead we want to join two data frames by multiple key variables. For example, we see that in order to join the flights and weather data frames, we need more than one key variable: year, month, day, hour, and origin. This is because the combination of these 5 variables act to uniquely identify each observational unit in the weather data frame: hourly weather recordings at each of the 3 NYC airports

```
flights_weather_joined <- flights %>%
  inner_join(weather, by = c("year", "month", "day", "hour", "origin"))

# explore it
glimpse(flights_weather_joined)
```

```
## Rows: 335,220
## Columns: 33
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013...
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 55...
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 60...
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2,...
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 8...
```

```
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 8...
## $ arr_delay      <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7,...
## $ carrier        <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6"...
## $ flight         <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301...
## $ tailnum        <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N...
## $ origin         <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LG...
## $ dest           <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IA...
## $ air_time       <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149...
## $ distance       <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 73...
## $ hour           <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6...
## $ minute         <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59...
## $ time_hour.x    <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-0...
## $ gain           <dbl> -9, -16, -31, 17, 19, -16, -24, 11, 5, -10, 0, 1, -9...
## $ hours          <dbl> 3.7833333, 3.7833333, 2.6666667, 3.0500000, 1.9333333...
## $ gain_per_hour  <dbl> -2.3788546, -4.2290749, -11.6250000, 5.5737705, 9.82...
## $ temp           <dbl> 39.02, 39.92, 39.02, 39.02, 39.92, 39.02, 37.94, 39....
## $ dewp           <dbl> 28.04, 24.98, 26.96, 26.96, 24.98, 28.04, 28.04, 24....
## $ humid          <dbl> 64.43, 54.81, 61.63, 61.63, 54.81, 64.43, 67.21, 54....
## $ wind_dir       <dbl> 260, 250, 260, 260, 260, 260, 240, 260, 260, 260, 26...
## $ wind_speed     <dbl> 12.65858, 14.96014, 14.96014, 14.96014, 16.11092, 12...
## $ wind_gust      <dbl> NA, 21.86482, NA, NA, 23.01560, NA, NA, 23.01560, NA...
## $ precip         <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...
## $ pressure       <dbl> 1011.9, 1011.4, 1012.1, 1012.1, 1011.7, 1011.9, 1012...
## $ visib          <dbl> 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, 10, ...
## $ time_hour.y    <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-0...
## $ temp_in_C      <dbl> 3.9, 4.4, 3.9, 3.9, 4.4, 3.9, 3.3, 4.4, 3.3, 4.4, 3....
```

```
head(flights_weather_joined)
```

```
## # A tibble: 6 x 33
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517           515           2       830           819
## 2  2013     1     1     533           529           4       850           830
## 3  2013     1     1     542           540           2       923           850
## 4  2013     1     1     544           545          -1      1004          1022
## 5  2013     1     1     554           600          -6       812           837
## 6  2013     1     1     554           558          -4       740           728
## # ... with 25 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour.x <dtm>, gain <dbl>, hours <dbl>,
## #   gain_per_hour <dbl>, temp <dbl>, dewp <dbl>, humid <dbl>, wind_dir <dbl>,
## #   wind_speed <dbl>, wind_gust <dbl>, precip <dbl>, pressure <dbl>,
## #   visib <dbl>, time_hour.y <dtm>, temp_in_C <dbl>
```

1.7 select variables/columns

flights data frame has 19 variables. Say you only need two of these 19 variables, say carrier and flight

```
just_two <- flights %>%
  select(carrier, flight)
```

```
head(just_two)
```

```
## # A tibble: 6 x 2
##   carrier flight
##   <chr>    <int>
## 1 UA      1545
## 2 UA      1714
## 3 AA      1141
## 4 B6       725
## 5 DL       461
## 6 UA      1696
```

Let's say instead you want to drop, or de-select, certain variables. For example, consider the variable `year` in the `flights` data frame. This variable isn't quite a "variable" because it is always 2013 and hence doesn't change. Remove this variable from the data frame by using the `"-"` sign

```
flights_no_year <- flights %>%
  select(-year)

head(flights_no_year)
```

```
## # A tibble: 6 x 21
##   month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1     1     1     517             515           2     830           819
## 2     1     1     533             529           4     850           830
## 3     1     1     542             540           2     923           850
## 4     1     1     544             545          -1    1004          1022
## 5     1     1     554             600          -6     812           837
## 6     1     1     554             558          -4     740           728
## # ... with 14 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, gain <dbl>, hours <dbl>,
## #   gain_per_hour <dbl>
```

Another way of selecting columns/variables is by specifying a range of columns

```
flight_arr_times <- flights %>%
  select(month:day, arr_time:sched_arr_time)

head(flight_arr_times)
```

```
## # A tibble: 6 x 4
##   month   day arr_time sched_arr_time
##   <int> <int>   <int>         <int>
## 1     1     1     830           819
## 2     1     1     850           830
## 3     1     1     923           850
## 4     1     1    1004          1022
## 5     1     1     812           837
## 6     1     1     740           728
```


The `select()` function can also be used to reorder columns when used with the `everything()` helper function. Suppose we want the `hour`, `minute`, and `time_hour` variables to appear immediately after the `year`, `month`, and `day` variables, while not discarding the rest of the variables

```
flights_reorder <- flights %>%
  select(year, month, day, hour, minute, time_hour, everything())

glimpse(flights_reorder)
```

```
## Rows: 336,776
## Columns: 22
## $ year      <int> 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013, 2013...
## $ month     <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ day       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ hour      <dbl> 5, 5, 5, 5, 6, 5, 6, 6, 6, 6, 6, 6, 6, 6, 5, 6, 6...
## $ minute    <dbl> 15, 29, 40, 45, 0, 58, 0, 0, 0, 0, 0, 0, 0, 0, 59...
## $ time_hour <dtm> 2013-01-01 05:00:00, 2013-01-01 05:00:00, 2013-01-0...
## $ dep_time  <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 55...
## $ sched_dep_time <int> 515, 529, 540, 545, 600, 558, 600, 600, 600, 600, 60...
## $ dep_delay <dbl> 2, 4, 2, -1, -6, -4, -5, -3, -3, -2, -2, -2, -2, -2,...
## $ arr_time  <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 8...
## $ sched_arr_time <int> 819, 830, 850, 1022, 837, 728, 854, 723, 846, 745, 8...
## $ arr_delay <dbl> 11, 20, 33, -18, -25, 12, 19, -14, -8, 8, -2, -3, 7,...
## $ carrier   <chr> "UA", "UA", "AA", "B6", "DL", "UA", "B6", "EV", "B6"...
## $ flight    <int> 1545, 1714, 1141, 725, 461, 1696, 507, 5708, 79, 301...
## $ tailnum   <chr> "N14228", "N24211", "N619AA", "N804JB", "N668DN", "N...
## $ origin    <chr> "EWR", "LGA", "JFK", "JFK", "LGA", "EWR", "EWR", "LG...
## $ dest      <chr> "IAH", "IAH", "MIA", "BQN", "ATL", "ORD", "FLL", "IA...
## $ air_time  <dbl> 227, 227, 160, 183, 116, 150, 158, 53, 140, 138, 149...
## $ distance  <dbl> 1400, 1416, 1089, 1576, 762, 719, 1065, 229, 944, 73...
## $ gain      <dbl> -9, -16, -31, 17, 19, -16, -24, 11, 5, -10, 0, 1, -9...
## $ hours     <dbl> 3.7833333, 3.7833333, 2.6666667, 3.0500000, 1.9333333...
## $ gain_per_hour <dbl> -2.3788546, -4.2290749, -11.6250000, 5.5737705, 9.82...
```

Lastly, the helper functions `starts_with()`, `ends_with()`, and `contains()` can be used to select variables/columns that match those conditions

```
a_var <- flights %>%
  select(starts_with("a"))

head(a_var)
```

```
## # A tibble: 6 x 3
##   arr_time arr_delay air_time
##   <int>     <dbl>     <dbl>
## 1     830         11        227
## 2     850         20        227
## 3     923         33        160
## 4    1004        -18        183
## 5     812        -25        116
## 6     740         12        150
```

```
delay_var <- flights %>%
  select(ends_with("delay"))

head(delay_var)
```

```
## # A tibble: 6 x 2
##   dep_delay arr_delay
##   <dbl>     <dbl>
## 1         2         11
## 2         4         20
## 3         2         33
## 4        -1        -18
## 5        -6        -25
## 6        -4         12
```

```
time_var <- flights %>%
  select(contains("time"))

head(time_var)
```

```
## # A tibble: 6 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
##   <int>         <int>     <int>         <int>     <dbl> <dtm>
## 1     517           515       830           819       227 2013-01-01 05:00:00
## 2     533           529       850           830       227 2013-01-01 05:00:00
## 3     542           540       923           850       160 2013-01-01 05:00:00
## 4     544           545      1004          1022       183 2013-01-01 05:00:00
## 5     554           600       812           837       116 2013-01-01 06:00:00
## 6     554           558       740           728       150 2013-01-01 05:00:00
```

1.8 rename variables

`rename` changes the name of variables. In `flights`, change `dep_time` and `arr_time` to be `departure_time` and `arrival_time`

```
flights_time_new <- flights %>%
  select(dep_time, arr_time) %>%
  rename(departure_time = dep_time, arrival_time = arr_time)

glimpse(flights_time_new)
```

```
## Rows: 336,776
## Columns: 2
## $ departure_time <int> 517, 533, 542, 544, 554, 554, 555, 557, 557, 558, 55...
## $ arrival_time    <int> 830, 850, 923, 1004, 812, 740, 913, 709, 838, 753, 8...
```

1.9 top_n values of a variable

return a data frame of the top 10 destination airports

```
top_10 <- named_dests %>%
  top_n(n = 10, wt = num_flights)

top_10
```

```
## # A tibble: 10 x 9
##   dest num_flights airport_name lat lon alt tz dst tzone
##   <chr>      <int> <chr>      <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1 ORD        17283 Chicago Ohare Intl 42.0 -87.9 668 -6 A America~
## 2 ATL        17215 Hartsfield Jackson~ 33.6 -84.4 1026 -5 A America~
## 3 LAX        16174 Los Angeles Intl 33.9 -118. 126 -8 A America~
## 4 BOS        15508 General Edward Law~ 42.4 -71.0 19 -5 A America~
## 5 MCO        14082 Orlando Intl 28.4 -81.3 96 -5 A America~
## 6 CLT        14064 Charlotte Douglas ~ 35.2 -80.9 748 -5 A America~
## 7 SFO        13331 San Francisco Intl 37.6 -122. 13 -8 A America~
## 8 FLL        12055 Fort Lauderdale Ho~ 26.1 -80.2 9 -5 A America~
## 9 MIA        11728 Miami Intl 25.8 -80.3 8 -5 A America~
## 10 DCA        9705 Ronald Reagan Wash~ 38.9 -77.0 15 -5 A America~
```